

Securing Applications in Kubernetes

Tsvi Korren, CISSP

April 2017



CONTAINERS HAVE SECURITY IMPLICATIONS



Risk posture of the images is not understood



Where should security fit in the process



Containers are not visible with current security tools

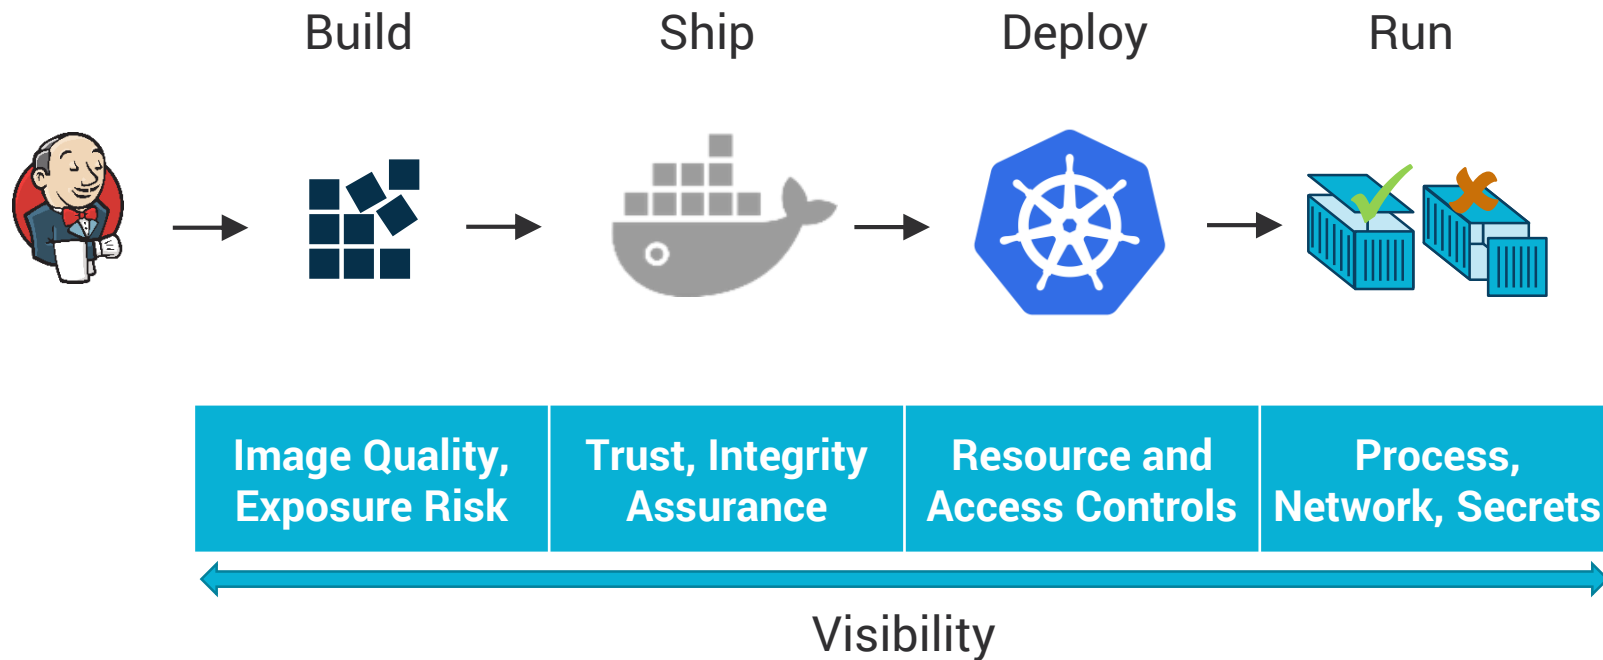


Open source and other external components used

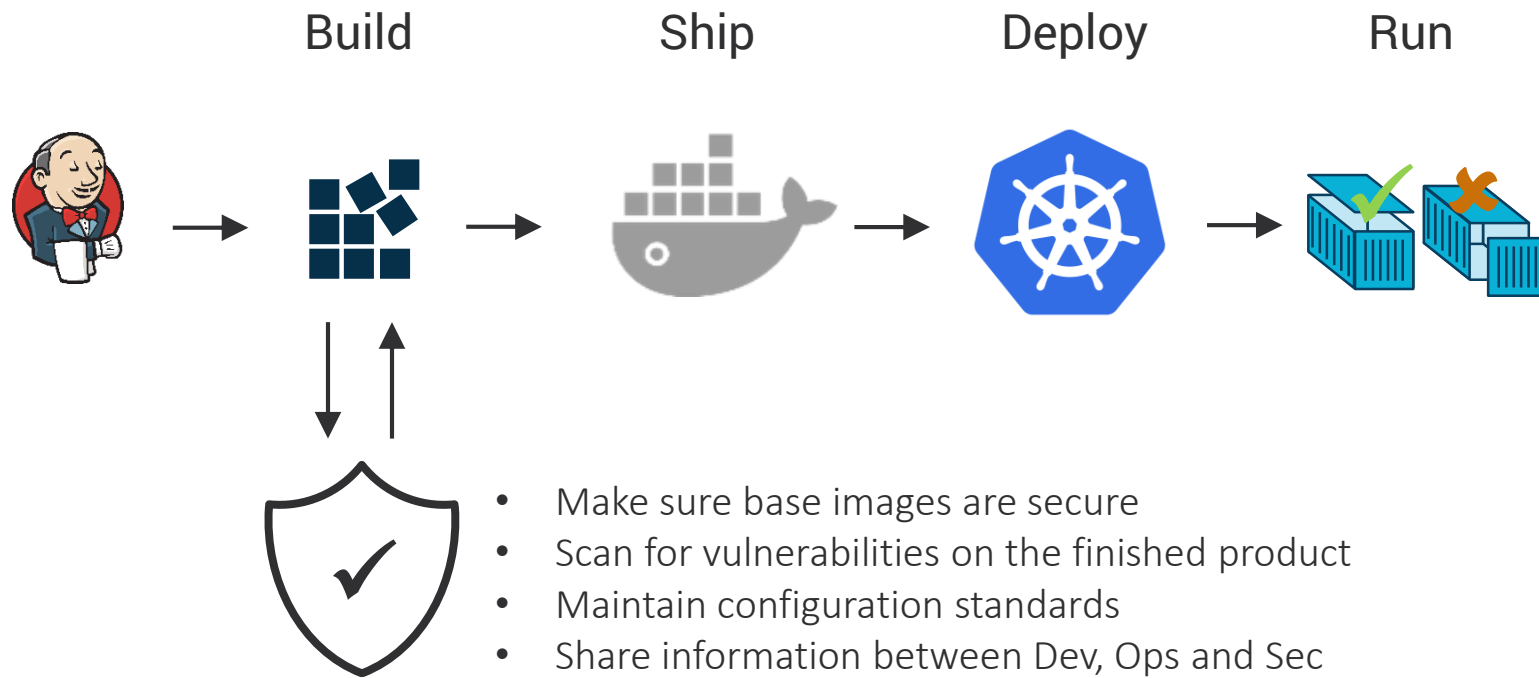
WHAT SECURITY PEOPLE WANT?

- Safe images, from trusted sources, tamper-proof
- Common security practices in the container environment
- Networking segmentation
- Safeguard sensitive data
- Accountability and audit data of container usage
- Data for demonstrating compliance

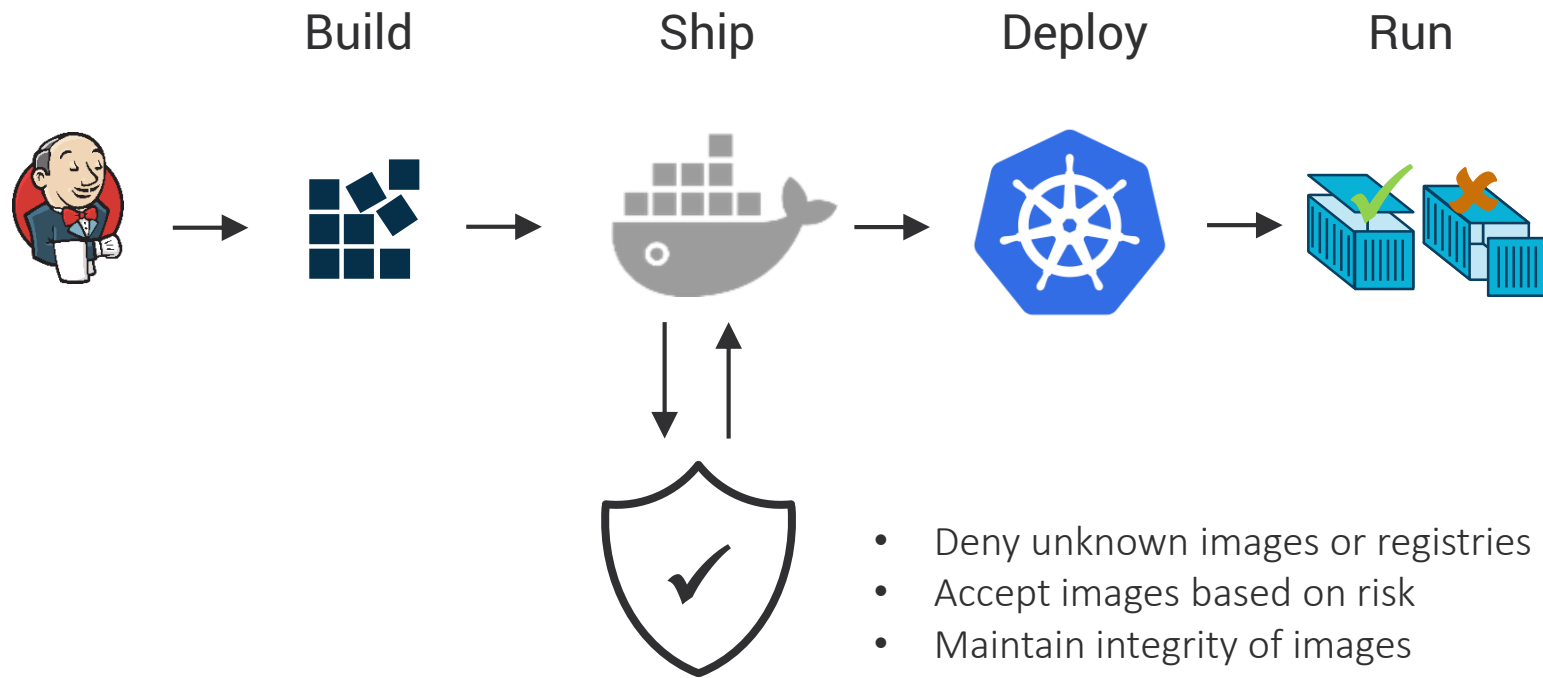
SECURE EACH STEP IN THE CONTAINER LIFECYCLE



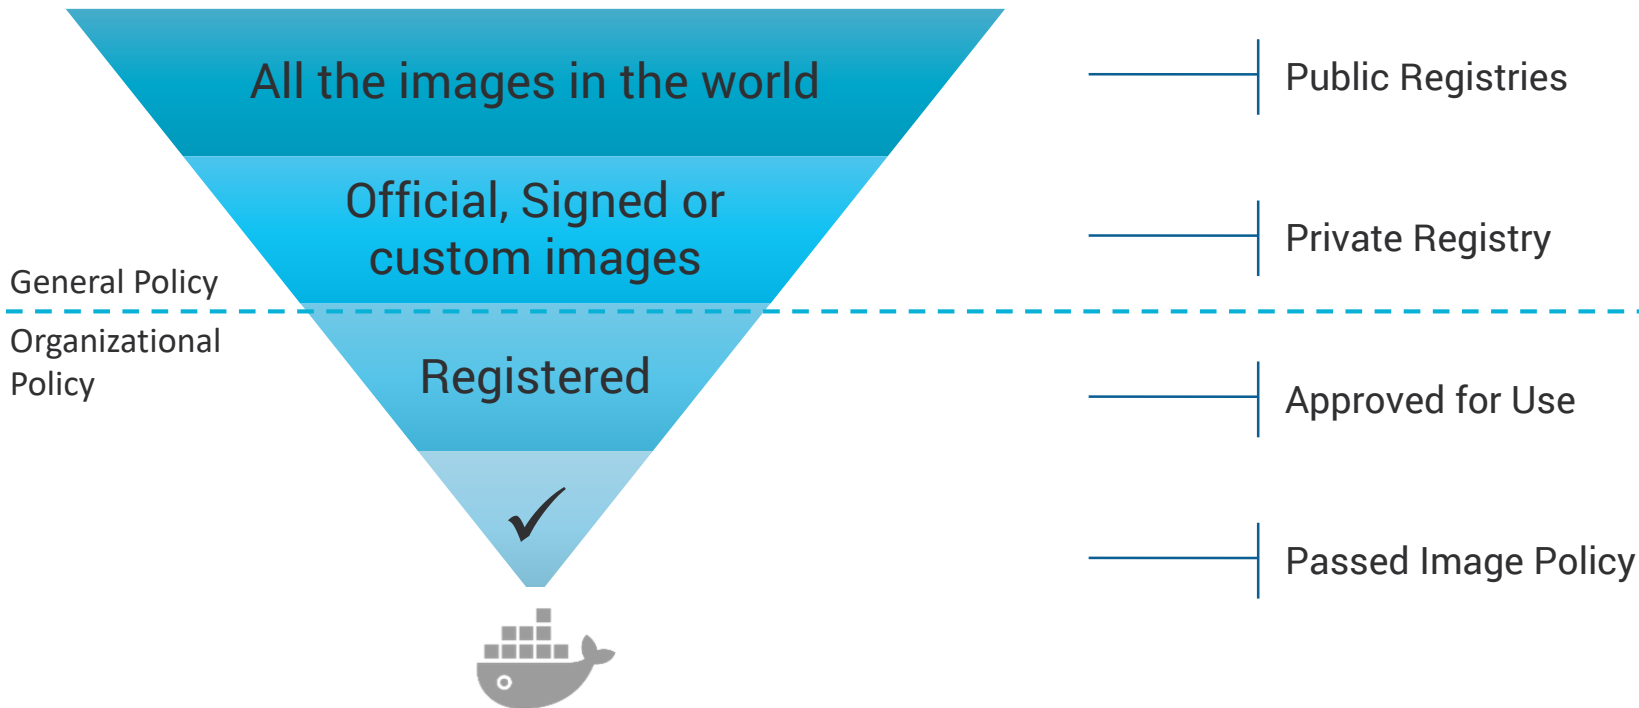
SECURITY STARTS IN THE BUILD PHASE



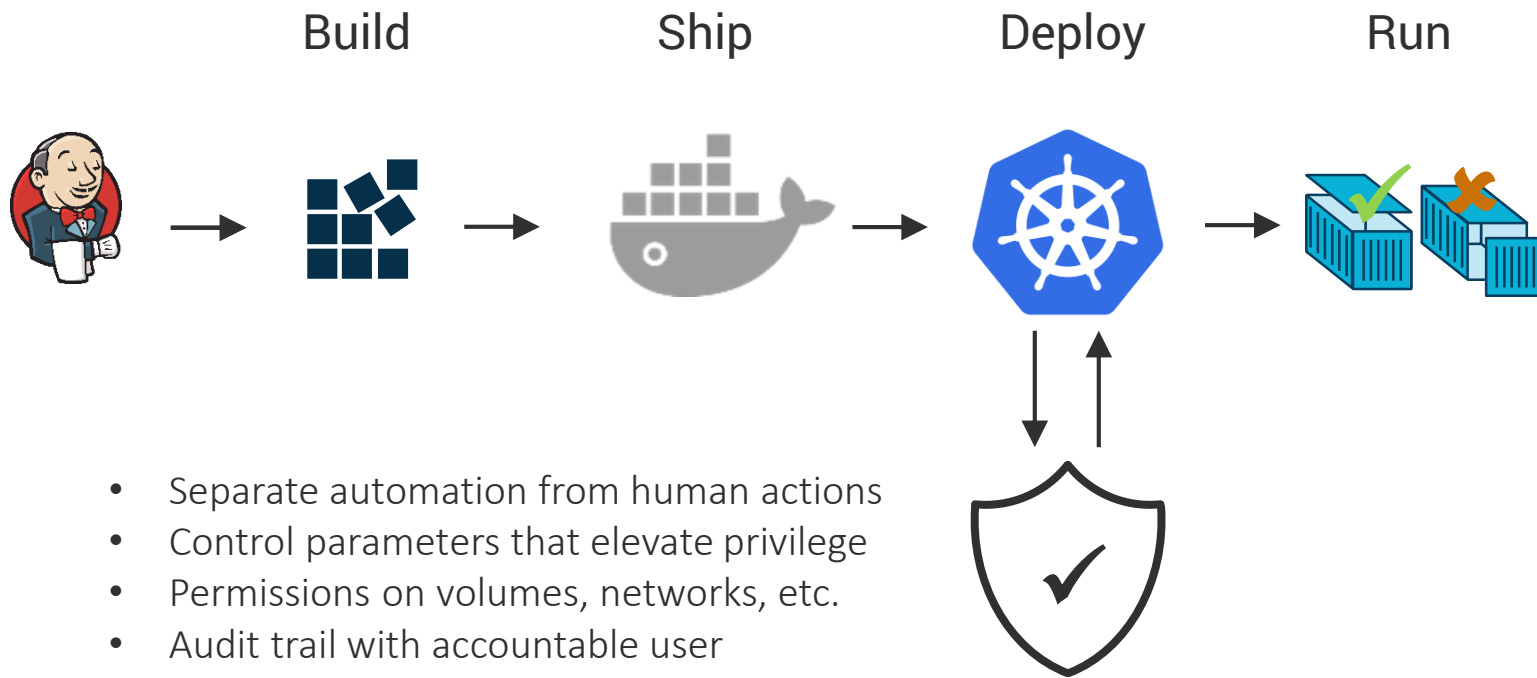
ADD ENFORCEMENT OF IMAGE USAGE



CONTROL THE INFLOW OF IMAGES



LIMIT ACCESS TO CONTAINER ENGINE



PROTECT NODES DOCKER ENGINE

- Limit terminal access to Kubernetes nodes
- Keep the Docker API secure
- Use kubectl with proper authorization
 - To manage containers via pods, deployments...
 - To exec into a running container, if needed...
 - To query status via describe...



KUBERNETES AUTHORIZATION

- Enables define fine-grained-access controls on
 - Namespaces
 - Pods, Services, Containers
 - Operations
- Authorization plugins based on
 - ABAC model (attribute-based)
 - RBAC mode (role-based)

ADMINISTRATIVE BOUNDARIES

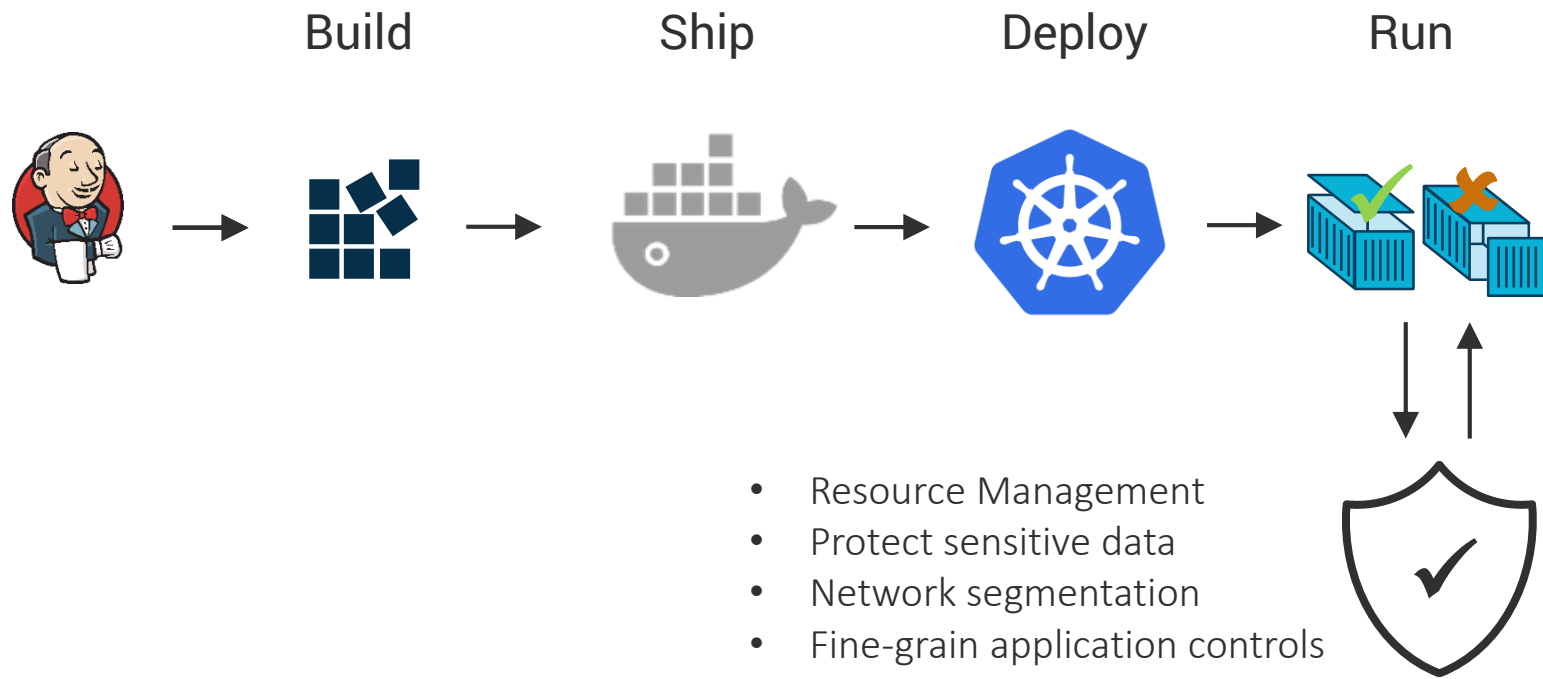
- Example: allow 'alice' to read pods from namespace 'fronto'

```
{
  "apiVersion": "abac.authorization.kubernetes.io/v1beta1",
  "kind": "Policy",
  "spec": {
    "user": "alice",
    "namespace": "fronto",
    "resource": "pods",
    "readonly": true
  }
}
```

WHY AUTHORIZATION?

- Limits the damage of mistake or malicious intent
- Manage resources access by logical groups
- Resource segregation for multi-tenancy
- Compliance with regulations

GRANULAR CONTROLS OF RUNNING CONTAINERS



DEFINE RESOURCE QUOTA

- Avoid Resource-unbound containers in shared cluster
- Create resource quota policies
 - Pods
 - CPUs
 - Memory ...
- Assigned to namespace

RESOURCE QUOTA EXAMPLE

■ compute-resources.yaml

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
spec:
  hard:
    pods: "4"
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
```

■ `kubectl create -f ./compute-resources.yaml --namespace=myspace`

WHY RESOURCE QUOTAS?

- Unbound resources add risk of DoS by runaway container
- By default, all resources are created unbound
- Predictable performance of clusters
- Required for capacity planning and disaster recovery

PROTECT SENSITIVE INFORMATION

- Storing sensitive data inside images or deployment definition is not safe
- Basic requirements for secret management
 - Put them in a vault while at rest, encrypted with key management
 - Restrict access to authorized users on authorized resources
 - Protect in transit and never commit to storage
 - Facilitate secrets expiry and rotation

KUBERNETES SECRETS

- Defined as objects consistent with the deployments
- Can be distributed as files or environment variables
- Cautions
 - Simple base64 encoding
 - Values accessible while at rest in etcd
 - No separation of duties: operator can see secret value
 - Secrets might persist on the node regardless of actual usage by containers

KUBERNETES SECRETS - EXAMPLE

```
$ echo -n "admin" | base64
YWRtaW4=
$ echo -n "1f2d1e2e67df" | base64
MWYyZDF1MmU2N2Rm
```

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDF1MmU2N2Rm
```

```
echo "MWYyZDF1MmU2N2Rm" | base64 --decode
1f2d1e2e67df
```

NETWORK SEGMENTATION

- Integration with external network enforcement points
- Kubernetes Network Policies work on pod-to-pod isolation (with only incoming traffic rules)
- Dynamic nature of container network identities makes container network segmentation a true challenge

IMPLEMENT NETWORK SEGMENTATION: EXAMPLE

```
apiVersion: extensions/v1beta1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  ingress:
    -
      from:
        -
          namespaceSelector:
            matchLabels:
              project: myproject
        -
          podSelector:
            matchLabels:
              role: frontend
      ports:
        -
          port: 6379
          protocol: tcp
  podSelector:
    matchLabels:
      role: db
```

WHY NETWORK SEGMENTATION

- One compromised application is an door open into the cluster
- Ensures containers communicate based on required function
- Enables more co-locating of applications in the cluster
- Network segmentation is required for compliance



BEYOND KUBERNETES NATIVE CONTROLS

- Image Assurance
- User Access controls
- Application granular controls
- Secrets distribution
- Network segmentation

DETAILED IMAGE RISK INFORMATION

Repositories & Images > centos:7

Vulnerabilities Packages Metadata History

Image Overview

6 High 12 Medium 1 Low 3.7 Average Score

CVE	SEVERITY	PACKAGE
> CVE-2016-5636	High	python
> CVE-2016-5636	High	python-libs
> CVE-2016-2834	High	nss-util
▼ CVE-2016-2834	High	nss

Description: Mozilla Network Security Services (NSS) before 3.23, as used in Mozilla Firefox before 47.0, allows remote attackers to cause unspecified other impact via unknown vectors.

CVSS v2 Score: 9.3

Vector: AV:N/AC:M/Au:N/C/C/I/C/A/C

Fix Version: nss-3.21.3-2.el7_3

NVD Reference: CVE-2016-2834

Vendor Reference: RHSA-2016:2779

For Ops and Security

For Developers

Jenkins

Jenkins > Peekr > #24 > Aqua Security Scanner

Back to Project
Status
Changes
Console Output
Edit Build Information
Delete Build
Aqua Security Scanner
Previous Build
Next Build

Vulnerability Report: peekr/demo:latest

From Registry: Docker Hub

1	2	0	5.7
HIGH	MEDIUM	LOW	SCORE AVG.

The following vulnerabilities were found:

Name	File	Severity	Score	Publish Date
CVE-2016-2515	/usr/share/nginx/html/js/utlis.js	high	7.8	2016-04-13
CVE-2016-2515	/usr/share/nginx/html/js/utlis.js	high	7.8	2016-04-13

PREVENT UNAUTHORIZED IMAGES FROM RUNNING

☒ Enable Package blacklist

List of blacklisted Packages:

Enter Package name: Version: Less Than: Enter Package Version: + Add

jre Less Than 1.7

Image CVEs Average Score Limit

If the average score of an image's CVEs equals to or exceeds this limit, the image will be disallowed automatically

☒ Enable average score limit

0 1 2 3 4 5 6 7 8 9 10

Image Policy

redis default 1/1 Docker Hub

☐ redis:latest (disallowed) ⓘ 5 High 13 Medium 3 Low

Image Assurance Policy
Action taken: Image **disallowed**
Reason: Maximum image CVE score **10.0** exceeded the maximum score limit **8.0**

Registration

```
core@ip-10-78-120-5 ~ $ docker pull redis:latest
Error response from daemon: You do not have permission to execute this command.
Use of unauthorized image.
core@ip-10-78-120-5 ~ $
```

Prevention

LIMIT HUMAN INTERACTION WITH AUTOMATION

User Access Control Policies [+ ADD NEW POLICY](#) [DELETE](#) [REFRESH](#)

10

<input type="checkbox"/>	NAME	ROLE	UPDATE TIME	AUTHOR
<input type="checkbox"/>	rule-admin-containers	Administrator	2016-11-08 08:31:35 PM	system
<input type="checkbox"/>	rule-admin-images	Administrator	2016-11-08 08:31:35 PM	system
<input type="checkbox"/>	rule-admin-networks	Administrator	2016-11-08 08:31:35 PM	system
<input type="checkbox"/>	rule-admin-swarm	Administrator	2016-11-08 08:31:35 PM	system
<input type="checkbox"/>	rule-admin-volumes	Administrator	2016-11-08 08:31:35 PM	system
<input type="checkbox"/>	rule-group-docker	Docker User	2016-11-08 08:31:35 PM	system
<input type="checkbox"/>	rule-power-users	Power User	2016-11-08 08:31:35 PM	system

Showing 1 to 7 of 7 entries

Previous **1** Next

```
alice@ip-10-78-120-5 / $ docker stop mongo
Error response from daemon: You do not have permission to execute this command. No matching rule granting access to resource
alice@ip-10-78-120-5 / $
```


APPLICATION GRANULAR CONTROLS

Resources			Network	Environment Variables	User Accounts
RESOURCE	ACCESS	TIME			
/usr/bin/bash	exec	2016-05-25 11:52:55 AM			
/usr/bin/dirname	exec	2016-05-25 11:52:55 AM			
/usr/bin/basename	exec	2016-05-25 11:52:55 AM			
/usr/bin/uname	exec	2016-05-25 11:52:55 AM			
/usr/bin/grep	exec	2016-05-25 11:52:55 AM			
/usr/lib/jvm/java/bin/java	exec	2016-05-25 11:52:55 AM			
/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.91-0.b14.el7 ...	exec	2016-05-25 11:52:55 AM			


Resources		Network	Environment Variables	User Accounts
UID	NAME			
1000	jboss			

```
secdemo-4 / # docker exec -it -u root app bash
Permission denied
secdemo-4 / # docker exec -it app sh
sh-4.2$ ping
sh: /usr/bin/ping: Permission denied
sh-4.2$ cp
sh: /usr/bin/cp: Permission denied
sh-4.2$ yum
sh: /usr/bin/yum: /usr/bin/python: bad interpreter
sh-4.2$
```

HOST NETWORK ZONES ARE NOT ENOUGH

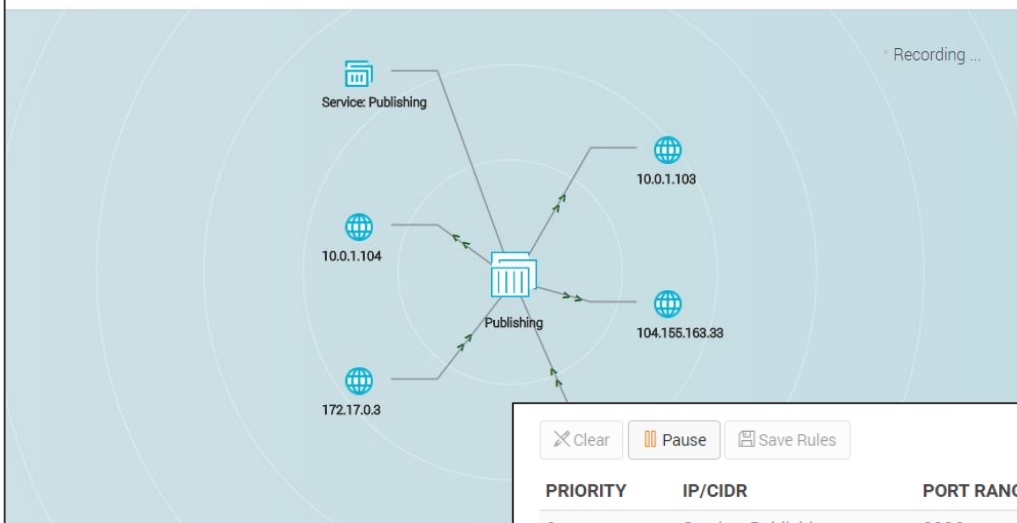


NEED FOR CONTAINER-SPECIFIC NETWORK ZONES



LEARN AND APPLY FIREWALL RULES

Services > Service Network for Publishing



Clear Pause Save Rules

PRIORITY	IP/CIDR	PORT RANGE	ALLOW/DENY	INBOUND / OUTBOUND	
0	Service: Publishing	3306	Allow	Outbound	
1	10.0.1.103	8069	Allow	Outbound	
2	10.0.1.104	81	Allow	Outbound	
3	104.155.163.33	80	Allow	Outbound	
4	172.17.0.3	3306	Allow	Inbound	

SECRETS MANAGEMENT

Secrets

Define and edit secrets that you plan to use in your container environment

Save Secret

Name	Value	Source	Description	Containers	Labels
db.password	*****	aqua		1	Select labels...

NAME ^	IMAGE	HOST	STATUS
app	demo:444/myapp:1.0	secdemo-4	▶ Running

Define
Secret

```
secdemo-4 / # docker run -d -e MYDB_ID=appdbuser -e MYDB_TOKEN=ToKeN -e MYDB_PWD={db.password} --name=app demo:444/myapp:1.0
dd94c492b55ee81af13dd7c590440c174d1839eabace28331fe3d3552d758f77
secdemo-4 / # docker inspect app | grep DB
    "MYDB_ID=appdbuser",
    "MYDB_PWD={db.password}",
    "MYDB_TOKEN=aqua-enc:7Ci9UEaZ1SE/sxiPxXT8iEBwiv5qz0oJKvKSTckukA0=",
secdemo-4 / # docker exec -it app bash -c set | grep DB
MYDB_ID=appdbuser
MYDB_PWD=MyNewValue
MYDB_TOKEN=ToKeN
```

True
Encryption

ADDED VISIBILITY

> 22 Jan 10:14:33 PM	Success	User root ran command <code>docker rm</code> on host devdemo-2
> 22 Jan 10:14:33 PM	Success	User root ran command <code>docker rm</code> on host devdemo-2
> 22 Jan 10:14:33 PM	Success	User root ran command <code>docker rm</code> on host devdemo-2
> 22 Jan 10:14:33 PM	Success	User root ran command <code>docker rm</code> on host devdemo-2
> 22 Jan 10:13:55 PM	Success	User administrator performed delete repository on <code>jboss/wildfly @Local Host</code>
> 22 Jan 10:13:16 PM	Success	User administrator performed delete repository on <code>aquasec/demo @Docker Hub</code>
> 22 Jan 10:12:53 PM	Success	User administrator performed update host on <code>devdemo-4</code>
> 22 Jan 10:12:26 PM	Success	User tsvi ran command <code>docker ps</code> on host devdemo-2
> 22 Jan 10:12:15 PM	Success	User tsvi ran command <code>docker ps</code> on host devdemo-2
> 22 Jan 10:12:10 PM	Block	User tsvi ran command <code>docker ps</code> on host devdemo-2
> 22 Jan 10:12:05 PM	Success	User tsvi ran command <code>docker ps</code> on host devdemo-2
> 22 Jan 10:12:05 PM	Success	User tsvi ran command <code>docker ps</code> on host devdemo-2
> 22 Jan 10:12:05 PM	Success	User tsvi ran command <code>docker ps</code> on host devdemo-2

Audit trail

Splunk
Syslog
OMS
(others)

Host Images


NAME	HOST	SECURITY STATUS	COMPLIES WITH POLICY	REGISTERED
> alpine:latest	tsvi-devdemo-2	✓ OK	✓ Yes	△ No
> odoo:9.0	tsvi-devdemo-2	○ Queued	-	△ No
> aquadev/scanner-cli:master	tsvi-devdemo-2	△ 2 High and 2 others...	✓ Yes	✓ Yes
> jenkins:docker	tsvi-devdemo-2	△ 30 High and 100 others...	✓ Yes	△ No
> myapp:latest	tsvi-devdemo-2	△ 7 High and 66 others...	✓ Yes	✓ Yes
> httpd:2.4.10	tsvi-devdemo-2	△ 48 High and 84 others...	✓ Yes	✓ Yes

Inventory

GRC

OPPORTUNITY FOR BETTER SECURITY

- Prevent unknown images
- Stop image by CVEs and score
- Stop user privilege escalation
- Stop suspicious processes
- Control capabilities
- Enforce network isolation
- Protect the host resources
- Encrypt sensitive variables
- Enforce use of automation tools
- Visibility across the environment



WWW.AQUASEC.COM

